

2. Parallel Programming Languages

- ²The main aim of parallel computing is to partition a computation into subcomputations that can be executed in parallel
- Several actions must occur in parallel programs, independently of how they are described
- The actions include:
 - ★ Data distribution among processes
 - ★ Work distribution among processes
 - ★ Data exchange among processes
 - ★ Synchronization of different processes
- Some of the identified distinct paths for the development of applications software for parallel computers include:

²M. Ganesh, MATH440/540, SPRING 2018

1. ³Extend an existing compiler to translate sequential programs into parallel programs
 2. Extend an existing language with new operations that allow users to express parallelism
 3. Add a new parallel language on top of an existing sequential language
 4. Define a parallel language and compiler system
- The first three approaches have been tried by a few companies and resulted in some products such as:
 - * Parallel Software Products (company)
 - * CODE (Computationally Oriented Display Environment)
 - * Hence (Heterogeneous Network Computing Environment)

³M. Ganesh, MATH440/540, SPRING 2018

- ⁴The main concern of the first three approaches is that the main focus is either on the language or on the compiler.
- Separate development of these two is likely to lead to difficulty in debugging codes with substantial complexity in application
- The fourth approach to create a parallel language and associated compiler systems gives the programmer the ability to express parallel operations explicitly
- One way for the fourth approach is to develop a parallel language from scratch.
- The programming language *OCCAM* is a well known example of this approach, which has a syntax strikingly different from traditional languages such as C and Fortran
- Another way in the fourth approach to support explicit parallelism is to add parallel constructs to an existing language
- *C** and Unified Parallel C (extensions of C) , and Fortran 90+ (variants of Fortran such as Fortran 95/2003 and HPF - High Performance Fortran) are examples of the fourth approach with extensions

⁴M. Ganesh, MATH440/540, SPRING 2018

- ⁵For example, one of the key extension of C in C^* is the concept called *shape*:
- A *shape* specifies the way in which parallel data are organized
- A *shape* is an array like object, but while array elements must be manipulated one at a time, elements of a shape may be manipulated simultaneously
- For example, to set up a template for parallel data/matrix of shape/size 100×200 , one may use the declaration
 - `shape [100][200]matrix;`
 - Then to create parallel data/matrix A, B, C , of size 100×200 , use
 - `float:matrix A, B, C;`
 - The statements that C is a sum of A and B , in C^* , is
 - `with (matrix) { C = A+B; }`
 - Thinking Machines Corporation developed C^*

⁵M. Ganesh, MATH440/540, SPRING 2018

- ⁶UPC (Unified Parallel C) is a relatively recent extension of the C programming language, designed for high performance computing on cluster parallel machines
- Commercial UPC compilers are developed, for example, by HP and IBM, see <http://h30097.www3.hp.com/upc/> and <http://www.alphaworks.ibm.com/tech/upccompiler>
- Open source UPC compilers are available, for example, GCC UPC, see <http://www.intrepid.com/upc.html>
- UPC is a SPMD (Single Program Multiple Data) extension of C
- The SPMD is a single program parallel approach that involves only one program which will execute differently on different processors
- The differences can occur through either implicit mechanism such as compiler directives or explicit constructs such as references to the processor number executing the code
- See the above websites for details of UPC facilitating explicit parallelism (such as MYTHREAD) with a shared address space (such as shared)

⁶M. Ganesh, MATH440/540, SPRING 2018

- ⁷Substantial extensions from Fortran 77 in Fortran 90+ (Fortran 95/2003, HPF) within the fourth approach facilitate parallelism
- For example, in Fortran 77 (or C), we need a double nested DO (or for) loops to add and/or subtract two $N \times M$ matrices A and B to get data C and/or D
- In Fortran 90+ this can be accomplished using FORALL statements
- The main difference between DO and FORALL in Fortran is that in the DO loop structure, execution must be in a strict order, while the statements in the FORALL construct may be executed in any order
- In particular, FORALL construct process the same set of elements in any order selected by the processor
- Fortran 95/2003 is an SPMD extension of Fortran 77/90
- HPF (High Performance Fortran) is a data parallel extension of Fortran 77/90
- The main focus of data parallelism is on distributing the data across different parallel computing nodes

⁷M. Ganesh, MATH440/540, SPRING 2018

⁸Most parallel programs are written using either

- MPI (Message Passing Interface) with a SPMD model:

- ★ Usually for scientific applications with C or Fortran (with wrappers used to call C from Fortran and vice-versa)
 - ★ Scales easily: user controlled data layout
 - ★ Difficulty: send/receive matching, message packing/unpacking

- or using shared memory programming model with OpenMP

- ★ Usually for non-scientific applications
 - ★ Easier to program: direct reads and writes to shared data
 - ★ Hard to scale: mostly limited to SMPs; no concept of locality

⁸M. Ganesh, MATH440/540, SPRING 2018

- ⁹Next we consider basic techniques in Fortran 95/2003,
- Then, we will develop parallel programming with MPI in F90+, C, and C++
- This will then be followed by parallel programming with OpenMP in F90+, C and C++
- Subsequently, we will look into MPMD models with codes written using F90, C, C++, and Python
- Time permitting, we will also explore parallel programming for GPU processors

⁹M. Ganesh, MATH440/540, SPRING 2018