

## 7. Parallel programming with MPI - II

## <sup>2</sup> More MPI details and commands

- For stopping a program (emergency) with MPI commands, use MPI\_Abort
- To abort with a failure code, say, fail\_code, use in Fortran, C, and C++ programming with MPI respectively:

```
CALL MPI_Abort(MPI_COMM_WORLD, fail_code, ierror)  
MPI_Abort(MPI_COMM_WORLD, fail_code);  
MPI::COMM_WORLD.Abort(fail_code);
```

- The state of MPI on a processing core can be tested using the following commands with started and stopped. In Fortran these two variables are logical and in C these are of the type int

```
CALL MPI_Initialized ( started , ierror )  
CALL MPI_Finalized ( stopped , ierror )  
ierror = MPI_Initialized ( & started );  
ierror = MPI_Finalized ( & stopped ) ;  
ierror = MPI::Is_initialized ( );  
ierror = MPI::Is_finalized ( ) ;
```

---

<sup>2</sup>M. Ganesh, MATH440/540, SPRING 2018

- <sup>3</sup>To synchronize/block so that all processing cores in a communicator wait until all cores have entered the call (and start up independently), use the MPI\_Barrier command, as below in Fortran, C, C++:

```
CALL MPI_Barrier(MPI_COMM_WORLD, ierror)
MPI_Barrier(MPI_COMM_WORLD);
MPI::COMM_WORLD.Barrier();
```

- To broadcast copies of the *same data*, say, buffer with datatype, say, DOUBLE\_PRECISION or DOUBLE, and total number of elements, say, count, to all processing cores in a communicator, say MPI\_COMM\_WORLD from, say, master use MPI\_Bcast command, as below in Fortran, C, C++:

```
CALL MPI_Bcast(buffer, count, MPI_DOUBLE_PRECISION, master, &
MPI_COMM_WORLD, ierror)
MPI_Bcast(buffer, count, MPI_DOUBLE, master, MPI_COMM_WORLD)
MPI::COMM_WORLD.Bcast(buffer, count, MPI::DOUBLE, master)
```

---

<sup>3</sup>M. Ganesh, MATH440/540, SPRING 2018

### Exercise 1:

Create a file with some data containing, say,  $N = 12$  numbers. Write a parallel (Fortran, C, or C++) program with MPI so that the master core reads the data from the file in, say, `my_array` and broadcasts `my_array` to all processing cores in a communicator. Your program should print `my_array` from each processing core and the associated core number.

### Exercise 2:

Take a copy of your Exercise 1 code. Modify the copy as follows: `my_array` should be read from the file instead by the third processing core in a communicator. (If the communicator contains only two cores, the program should print an error message and abort.) The third processing core should broadcast only the middle six elements of `my_array`.

Your modified program should print the six elements of `my_array` from each processing core and the associated core number.

- <sup>4</sup>To scatter copies of different data from, say, master to all processing cores in a communicator, by splitting a bigger data evenly between all processing cores in a communicator, use MPI\_SCATTER:
- Let the total size of communicator be given by num\_cores (That is, using previous notation, my\_rank = 0, 1, ..., num\_cores-1.)
- Suppose that we want master to send equal parts of a data, say, buffer with data type DOUBLE\_PRECISION or DOUBLE and total number of elements count = num\_cores \* send\_count
- There are num\_cores equal sized blocks of data in buffer with each block consisting of send\_count number of elements
- Suppose that we want the master to scatter buffer so that, for  $i = 1, \dots, \text{num\_cores}$ , the  $i$ -th processing core gets the  $i$ -th block of send\_count number of elements from buffer, as a variable recv\_buffer with its count recv\_count
- The send process is used only on master and the receive process is used on all processing cores
- To achieve the above task, use the following (F, C, C++) commands:

---

<sup>4</sup>M. Ganesh, MATH440/540, SPRING 2018

```
5CALL MPI_Scatter( &
buffer, send_count, MPI_DOUBLE_PRECISION, &
recv_buffer, recv_count, MPI_DOUBLE_PRECISION, &
master, MPI_COMM_WORLD, ierror)

MPI_Scatter( buffer, send_count, MPI_DOUBLE, recv_buffer, recv_count,
MPI_DOUBLE, master, MPI_COMM_WORLD)

MPI::COMM_WORLD.Scatter( buffer, send_count, MPI::DOUBLE, recv_buffer,
recv_count, MPI::DOUBLE, master)
```

- Note that `buffer` and `recv_buffer` may have different ranks. For example, `buffer` may be a matrix and `recv_buffer` may be a vector

---

<sup>5</sup>M. Ganesh, MATH440/540, SPRING 2018

### Exercise 3:

Create a file with some data containing, say,  $N = 12$  numbers. Write a parallel (Fortran, C, or C++) program with MPI so that the master core reads the data from the file in a  $3 \times 4$  matrix, say, `my_mat`.

Let `num_cores` be the total number of processing core in a communicator, with `master = 0` being the first processing core in the communicator.

Let `send_count = N/4`. The master core should scatter `my_mat` so that: the first processing core receives the first column of `my_mat`; the second processing core receives the second column of `my_mat` and so on. The processing cores should receive these elements in a new array `recv_array`.

Your program should check if `send_count*num_cores == N`. If this condition is violated, the program should abort.

Your program should be such that the master prints `my_mat` and its rank and each processing core prints its rank and `recv_array`.

- <sup>6</sup>MPI\_Gather command is the converse of the operation MPI\_Scatter
- For MPI\_Gather, the recv\_buffer is larger with size recv\_count\*num\_cores
- In Fortran, C, and C++ programming with MPI, respectively, use:

```
CALL MPI_Gather( &
send_buffer, send_count, MPI_DOUBLE_PRECISION, &
recv_buffer, recv_count, MPI_DOUBLE_PRECISION, &
master, MPI_COMM_WORLD, ierror)
```

```
MPI_Gather( send_buffer, send_count, MPI_DOUBLE, recv_buffer, recv_count,
MPI_DOUBLE, master, MPI_COMM_WORLD)
```

```
MPI::COMM_WORLD.Gather( send_buffer, send_count, MPI::DOUBLE, recv_buffer,
recv_count, MPI::DOUBLE, master)
```

---

<sup>6</sup>M. Ganesh, MATH440/540, SPRING 2018

- For example, to implement Exercise 4 in Topic 7, we may avoid MPI\_Send and MPI\_Recv and instead use

```
IF (my_rank == master) THEN  
    ALLOCATE(all_sum(0:num_cores-1))  
END IF
```

```
call MPI_Gather(local_sum, 1, MPI_DOUBLE_PRECISION, all_sum, 1, &  
MPI_DOUBLE_PRECISION, master, MPI_COMM_WORLD,ierror)
```

```
full_sum = sum(all_sum)
```

#### Exercise 4:

Generalize your program in Exercise 3 (current section) so that the each processing core computes the square root of each element in recv\_array and the master gathers all these results in a new  $3 \times 4$  matrix, say, my\_mat\_comp\_sqrt Your program should be such that the master prints my\_mat, and my\_mat\_comp\_sqrt and each processing core prints its rank and recv\_array.

- <sup>7</sup>We may gather the data on master and then broadcast it to all processing cores
- Instead this can be achieved using a single command MPI\_Allgather
- In this symmetric operation, the command is similar to MPI\_Gather except that we need to remove the master argument

```
CALL MPI_Allgather( &
send_buffer, send_count, MPI_DOUBLE_PRECISION, &
recv_buffer, recv_count, MPI_DOUBLE_PRECISION, &
MPI_COMM_WORLD, ierror)
```

```
MPI_Allgather(send_buffer, send_count, MPI_DOUBLE, recv_buffer, recv_count,
MPI_DOUBLE, MPI_COMM_WORLD)
```

```
MPI::COMM_WORLD.Allgather(send_buffer, send_count, MPI::DOUBLE, recv_buffer
recv_count, MPI::DOUBLE)
```

---

<sup>7</sup>M. Ganesh, MATH440/540, SPRING 2018

<sup>8</sup>**Exercise 5:**

Generalize your program in Exercise 3 so that the each processing core computes the square root of each element in `recv_array` and each processing core receives the results in a  $3 \times 4$  matrix, say, `my_mat_comp_sqrt`. Your program should be such that the master prints `my_mat`, and `my_mat_comp_sqrt` and each processing core prints its rank and `my_mat_comp_sqrt`.

---

<sup>8</sup>M. Ganesh, MATH440/540, SPRING 2018

- <sup>9</sup>We can achieve a combination of the scatter and gather operation using MPI\_Alltoall
- Use this if you want all processing cores to act like the master to scatter and gather data
- Such a powerful operation is useful for achieving, for example, a parallel transpose of a matrix:

```
CALL MPI_Alltoall( &
send_buffer, send_count, MPI_DOUBLE_PRECISION, &
recv_buffer, recv_count, MPI_DOUBLE_PRECISION, &
MPI_COMM_WORLD, ierror)
```

```
MPI_Alltoall(send_buffer, send_count, MPI_DOUBLE, recv_buffer, recv_count,
MPI_DOUBLE, MPI_COMM_WORLD)
```

```
MPI::COMM_WORLD.Alltoall(send_buffer, send_count, MPI::DOUBLE, recv_buffer,
recv_count, MPI::DOUBLE)
```

---

<sup>9</sup>M. Ganesh, MATH440/540, SPRING 2018

<sup>10</sup>**Exercise 6:**

Let  $P$  be the size (total number of cores) in a communicator.

Let  $A_P$  be a  $P \times P$  matrix defined by

$$A_P = \begin{bmatrix} 1 & 2 & \cdots & P \\ P+1 & P+2 & \cdots & 2P \\ 2P+1 & 2P+2 & \cdots & 3P \\ \cdots & \cdots & \cdots & \cdots \\ \cdots & \cdots & \cdots & \cdots \\ \cdots & \cdots & \cdots & \cdots \\ (P-1)P+1 & (P-1)P+2 & \cdots & P^2 \end{bmatrix}. \quad (7.1)$$

Write a parallel (Fortran, or C, or C++) program with MPI using  $P$  cores (determined by `MPI_COMM_SIZE`) so that for  $i = 1, \dots, P$ :

- The  $i$ -th processing core allocates a  $P$  dimensional vector `arow` with elements being the  $i$ -th row of  $A_P$
- Using the `arow` vectors and `MPI_Alltoall` command, the  $i$ -th processing core also gets the  $i$ -th row of  $A_P^T$
- The master prints the matrix  $A_P$  in a file `A_mat.txt` and the matrix  $A_P^T$  in a file `transp_A_mat.txt`. (Check these two files to validate code.)

---

<sup>10</sup>M. Ganesh, MATH440/540, SPRING 2018

- <sup>11</sup>If the purpose of the MPI\_Gather command in your MPI code is to achieve some reduction (such as summation, find maximum of values, or ...) of the gathered data using master, the reduction can be achieved using MPI\_Reduce operations, without the need to use MPI\_Gather:
- Unlike in MPI\_Scatter and MPI\_Gather commands, for reduce operations, it is sufficient to specify the datatype and count of the buffer only once
- It is important to note that the reduction process does not reduce over the vector (that is if the count of the buffer is more than one)
- If count is more than one, then the reduction operation is applied for each index in buffer separately
- In case of requiring reduction over a vector, you need to get this done by yourself first
- Some of the MPI operations are:  
MPI\_SUM, MPI\_MIN, MPI\_MAX, MPI\_PROD, MPI\_MAXLOC, MPI\_MINLOC, ....
- Below, we use oper for any one of these reduction operations

---

<sup>11</sup>M. Ganesh, MATH440/540, SPRING 2018

```
12CALL MPI_Reduce( &  
send_buffer, recv_buffer, count, MPI_DOUBLE_PRECISION, &  
oper, master, MPI_COMM_WORLD, ierror)  
  
MPI_Reduce(send_buffer, recv_buffer, count, MPI_DOUBLE, oper, master,  
MPI_COMM_WORLD)  
  
MPI::COMM_WORLD.Reduce(send_buffer, recv_buffer, count, MPI::DOUBLE, oper,  
master)
```

- For example, to implement Exercise 4 in Topic 7, we may avoid MPI\_Send and MPI\_Recv, or MPI\_Gather and instead use just one main MPI command:

```
CALL MPI_Reduce(local_sum, full_sum, 1, MPI_DOUBLE_PRECISION, &  
MPI_SUM, master, MPI_COMM_WORLD, ierror)
```

---

<sup>12</sup>M. Ganesh, MATH440/540, SPRING 2018

- <sup>13</sup> **Exercise 7:**

Generalize your program in Exercise 3 (current section) so that the master deallocates `my_mat` and then receives the maximum element in each column of `my_mat` from all processing cores. Then the master computes the maximum value in the destroyed `my_mat`.

Your program should be such that the master prints `my_mat` and the maximum element in `my_mat`.

- We may be interested in all processing cores (instead of just the master) to have the reduced result
- We can achieve this using a single command `MPI_Allreduce`
- In this symmetric operation, the command is similar to `MPI_Reduce` except that we need to remove the master argument

---

<sup>13</sup>M. Ganesh, MATH440/540, SPRING 2018

<sup>14</sup>CALL MPI\_Allreduce( &  
send\_buffer, recv\_buffer, count, MPI\_DOUBLE\_PRECISION, &  
oper, MPI\_COMM\_WORLD, ierror)

MPI\_Allreduce(send\_buffer, recv\_buffer, count, MPI\_DOUBLE, oper,  
MPI\_COMM\_WORLD)

MPI::COMM\_WORLD.Allreduce(send\_buffer, recv\_buffer, count, MPI::DOUBLE,  
oper)

### Exercise 8:

Generalize your program in Exercise 3 (current section) so that the master deallocates my\_mat. Then all processing cores compute the maximum value in the destroyed my\_mat.

Your program should be such that all processing core should print its rank and the maximum element in my\_mat.

---

<sup>14</sup>M. Ganesh, MATH440/540, SPRING 2018